

Memory Management

Lecture 9

LRU Page Replacement Algorithm

Minakshi R.

Operating System Concepts 8th edition silberschatz Galvin

Least Recently Used (LRU) Algorithm

- Use past knowledge rather than future
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page
- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used
- But how to implement?

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

LRU Algorithm (Cont.)

□ Counter implementation

- Every page entry has a counter; every time page is referenced through this entry, **copy the clock into the counter**
- When a page needs to be changed, look at the counters to find smallest value
 - Search through table needed

□ Stack implementation


- Top is most recently used. Bottom is least recently used.
- Since stack entries are removed from the middle of the stack, a **doubly-linked list** implementation is best
- (**6 pointer changes at worst**--2 for internal neighbors, top of stack, previous top's pointer, entry's 2 pointers)

- LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly


Use Of A Stack to Record The Most Recent Page References

reference string


4 7 0 7 1 0 1 2 1 2 7 1 2



stack
before
a



stack
after
b



LRU Approximation Algorithms

□ Implementation of LRU can be expensive, so in practice may approximate LRU

□ Reference bit

- every entry in page table has reference bit
- bit set (to **1**) **when page accessed**
- **periodically the OS “cleans”** the reference bits (resets all to 0)
- pages with bit reset (i.e., **0**) **have higher priority** to be replaced
- Reference bits tell us what pages have been accessed but **do not tell us the order of use**

□ Additional reference bits algorithm

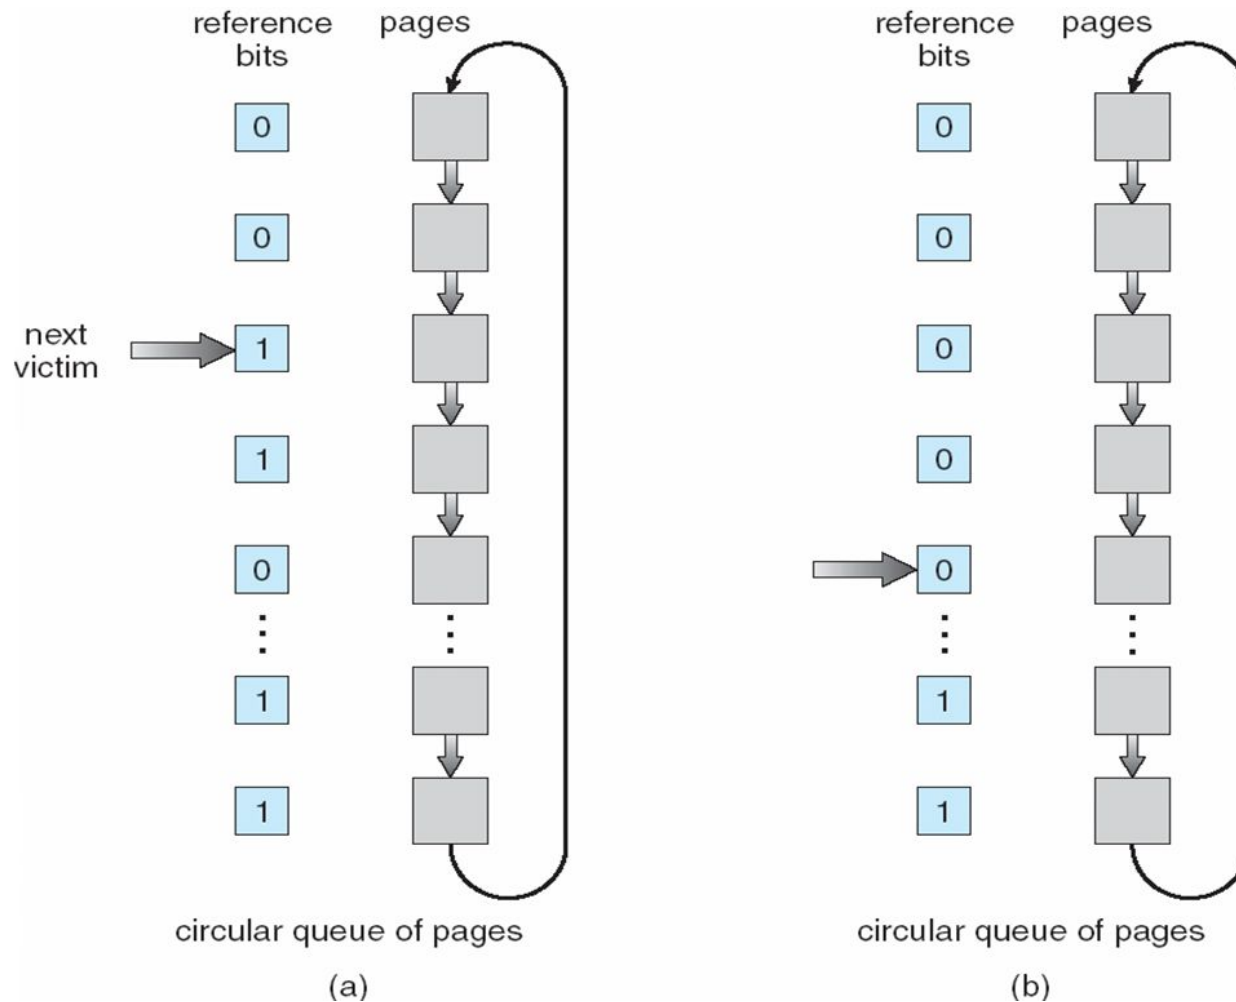
- keep record of reference bits at regular time intervals
- associate **8 bit byte** with each page table entry
- at regular intervals (e.g., 100 milliseconds) **OS shifts reference bit** for each page into high-order bit of 8-bit byte, shifting other bits right by one.
- low-order bit discarded
 - 00000000 not used in 8 time units
 - 11111111 used in each of past 8 time units
 - 00111111 not used in last 2 time units but used in each of 6 previous ones
- If view numbers as unsigned integers, the **page(s) with the lowest number is the one least recently used** (not guaranteed to be one page--can be ambiguous then)

LRU Approximation Algorithms

□ **Second-chance algorithm (aka clock algorithm)**

- FIFO replacement but inspect reference bit before actually replacing
- if reference bit == **0 then replace**
- if reference bit == **1 then clear reference bit and move on to inspect the next FIFO page**
- **Implementation: circular queue**; advance pointer, **clearing reference bits until find one with reference bit == 0**
- Worst case: cycle through all entries (degenerates to FIFO in this case)
- Slowly sweeping clock hand is good---means there is plenty of memory, not many page faults

Second-Chance (clock) Page-Replacement Algorithm




Counting Algorithms

- Keep a counter of the number of references that have been made to each page
 - Not common
- **LFU Algorithm:** replaces page with smallest count
- **MFU Algorithm:** based on the argument that the page with the smallest count was probably just brought in and has yet to be used

Page-Buffering Algorithms

- Keep a **pool of free frames**, always
 - Then frame available when needed, not found at fault time
 - Read page into free frame and select victim to evict and add to free pool
 - When convenient, evict victim
- Possibly, **keep list of modified pages**
 - When backing store otherwise idle, write pages there and set to non-dirty
- Possibly, keep free frame contents intact and note what is in them
 - If referenced again before reused, no need to load contents again from disk
 - Generally useful to reduce penalty if wrong victim frame selected



F= page fault occurring after the frame allocation is initially filled

Figure 8.15 Behavior of Four Page Replacement Algorithms

Comparison




Figure 8.17 Comparison of Fixed-Allocation, Local Page Replacement Algorithms

Applications and Page Replacement

- All of these algorithms have OS guessing about future page access
- Some applications have better knowledge – i.e. databases
- Memory intensive applications can cause double buffering
 - OS keeps copy of page in memory as I/O buffer
 - Application keeps page in memory for its own work
- Operating system can given direct access to the disk, getting out of the way of the applications
 - **Raw disk** mode
- Bypasses buffering, locking, etc

THANK YOU